


# Đề tài: Xử lý dữ liệu lớn với kỹ thuật Tách nền – Nhóm 2 –

## 64HTTT1

### Phần 1: Chạy mapreduce với Image Encoder

#### 1. Chuẩn bị dữ liệu và đẩy lên HDFS

Bước 1: Chuẩn bị ảnh và nén thành các file .tar



Bước 2: Tạo thư mục chứa các file .tar trên hdfs

**hdfs dfs -mkdir -p /hadoop/image\_shards**

The screenshot shows the HDFS Web UI with the URL `/hadoop`. The page title is "Browse Directory". The table header includes columns for Permission, Owner, Group, Size, Last Modified, Replication, Block Size, and Name. The data table shows the following entries:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	lonovo	supergroup	0 B	Dec 27 21:08	0	0 B	image_shards
drwxr-xr-x	lonovo	supergroup	0 B	Dec 27 18:40	0	0 B	input_shard_manifest
drwxr-xr-x	lonovo	supergroup	0 B	Dec 27 22:33	0	0 B	reports
drwxr-xr-x	lonovo	supergroup	0 B	Dec 27 21:54	0	0 B	zim_feature_output
drwxr-xr-x	lonovo	supergroup	0 B	Dec 27 22:32	0	0 B	zim_results_npz

Showing 1 to 5 of 5 entries

Bước 3: Đẩy file .tar lên hdfs

**hdfs dfs -put**

"D:\Hoc\_ki\_7\Phan\_tich\_du\_lieu\_lon\Code\ZIM\image\_shards\\*.tar"  
**/hadoop/image\_shards**

The screenshot shows the HDFS Web UI with the URL `/hadoop/image_shards`. The page title is "Browse Directory". The table header includes columns for Permission, Owner, Group, Size, Last Modified, Replication, Block Size, and Name. The data table shows the following entries:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-f--f--	lonovo	supergroup	3.08 MB	Dec 27 21:08	1	128 MB	shard_01.tar
-rw-f--f--	lonovo	supergroup	2.64 MB	Dec 27 21:08	1	128 MB	shard_02.tar
-rw-f--f--	lonovo	supergroup	2.44 MB	Dec 27 21:08	1	128 MB	shard_03.tar
-rw-f--f--	lonovo	supergroup	3.11 MB	Dec 27 21:08	1	128 MB	shard_04.tar
-rw-f--f--	lonovo	supergroup	2.76 MB	Dec 27 21:08	1	128 MB	shard_05.tar
-rw-f--f--	lonovo	supergroup	3.18 MB	Dec 27 21:08	1	128 MB	shard_06.tar
-rw-f--f--	lonovo	supergroup	2.76 MB	Dec 27 21:08	1	128 MB	shard_07.tar
-rw-f--f--	lonovo	supergroup	2.48 MB	Dec 27 21:08	1	128 MB	shard_08.tar
-rw-f--f--	lonovo	supergroup	2.36 MB	Dec 27 21:08	1	128 MB	shard_09.tar
-rw-f--f--	lonovo	supergroup	3.39 MB	Dec 27 21:08	1	128 MB	shard_10.tar

Bước 4: Tạo thư mục chứa file đường dẫn

**hdfs dfs -mkdir -p /hadoop/input\_shard\_manifest**

## Browse Directory

The screenshot shows a file browser interface for HDFS. The address bar at the top contains the path '/hadoop'. Below the address bar is a search bar labeled 'Search' and a set of icons for file operations like copy, move, and delete. A table below lists five entries in the directory:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	lonovo	supergroup	0 B	Dec 27 21:08	0	0 B	image_shards
drwxr-xr-x	lonovo	supergroup	0 B	Dec 27 18:40	0	0 B	input_shard_manifest
drwxr-xr-x	lonovo	supergroup	0 B	Dec 27 22:33	0	0 B	reports
drwxr-xr-x	lonovo	supergroup	0 B	Dec 27 21:54	0	0 B	zim_feature_output
drwxr-xr-x	lonovo	supergroup	0 B	Dec 27 22:32	0	0 B	zim_results_npz

At the bottom left, it says 'Showing 1 to 5 of 5 entries'. At the bottom right, there are buttons for 'Previous', '1' (highlighted), and 'Next'.

Bước 5: Đẩy file chứa đường dẫn lên hdfs

**hdfs dfs -put "**

**"D:\Hoc\_ki\_7\Phan\_tich\_du\_lieu\_lon\Code\ZIM\mapreduce\manifest\_shards.txt**  
**" /hadoop/input\_shard\_manifest**

The screenshot shows a code editor window with the file 'manifest\_shards.txt' open. The file contains a list of 20 paths, each starting with '/hadoop/image\_shards/shard\_'. The paths are numbered from 1 to 20 on the left side of the editor.

```
1 /hadoop/image_shards/shard_01.tar
2 /hadoop/image_shards/shard_02.tar
3 /hadoop/image_shards/shard_03.tar
4 /hadoop/image_shards/shard_04.tar
5 /hadoop/image_shards/shard_05.tar
6 /hadoop/image_shards/shard_06.tar
7 /hadoop/image_shards/shard_07.tar
8 /hadoop/image_shards/shard_08.tar
9 /hadoop/image_shards/shard_09.tar
10 /hadoop/image_shards/shard_10.tar
11 /hadoop/image_shards/shard_11.tar
12 /hadoop/image_shards/shard_12.tar
13 /hadoop/image_shards/shard_13.tar
14 /hadoop/image_shards/shard_14.tar
15 /hadoop/image_shards/shard_15.tar
16 /hadoop/image_shards/shard_16.tar
17 /hadoop/image_shards/shard_17.tar
18 /hadoop/image_shards/shard_18.tar
19 /hadoop/image_shards/shard_19.tar
20 /hadoop/image_shards/shard_20.tar
```

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

## Browse Directory

/hadoop/input\_shard\_manifest

Show 25 entries Search:


Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	lenovo	supergroup	698 B	Dec 27 18:40	1	128 MB	manifest_shards.txt

Showing 1 to 1 of 1 entries Previous 1 Next

Hadoop, 2020.

## 2. Chuẩn bị model và code để chạy mapreduce

Bước 1: Tải model encoder từ link: [https://huggingface.co/naver-iv/zim-anything-vitb/tree/main/zim\\_vit\\_b\\_2043](https://huggingface.co/naver-iv/zim-anything-vitb/tree/main/zim_vit_b_2043)



Bước 2: Xây dựng code Image\_Encoder có tên là zim\_encoder.py

```
#!/usr/bin/env python3
from typing import Any, Tuple

import numpy as np
import torch
import onnxruntime

def np2tensor(np_array: np.ndarray, device: torch.device) -> torch.Tensor:
    """Convert numpy array -> torch.Tensor trên device chỉ định."""
    return torch.from_numpy(np_array).to(device)

def tensor2np(torch_tensor: torch.Tensor) -> np.ndarray:
```

```

"""Convert torch.Tensor -> numpy array (luôn đưa về CPU trước)."""
return torch_tensor.detach().cpu().numpy()

class ZIM_Encoder:
    def __init__(
        self,
        onnx_path: str,
        num_threads: int = 8,
        use_cuda: bool = False,
        device_id: int = 0,
        verbose: bool = True,
    ):
        """
        onnx_path: đường dẫn tới file .onnx của image encoder ZIM
        num_threads: số thread cho onnxruntime
        use_cuda: nếu True sẽ cố gắng dùng CUDAExecutionProvider
        device_id: GPU id (nếu dùng CUDA)
        verbose: in thông tin input/output của model
        """
        self.onnx_path = onnx_path

        session_options = onnxruntime.SessionOptions()
        session_options.intra_op_num_threads = num_threads
        session_options.inter_op_num_threads = num_threads

        # Mặc định luôn có CPU provider
        providers = ["CPUExecutionProvider"]
        provider_options = []

        if use_cuda:
            # Thử thêm CUDA provider (nếu đã cài onnxruntime-gpu + có CUDA)
            try:
                providers.insert(0, "CUDAExecutionProvider")
                provider_options.insert(0, {"device_id": device_id})
                if verbose:
                    print(f"[ZIM_Encoder] Trying CUDAExecutionProvider on device {device_id}")
            except Exception as e:

```

```

        if verbose:
            print(f"[ZIM_Encoder] Cannot enable CUDAExecutionProvider,
fallback to CPU. Reason: {e}")

# Khởi tạo session
self.ort_session = onnxruntime.InferenceSession(
    onnx_path,
    sess_options=session_options,
    providers=providers,
    provider_options=provider_options if provider_options else None,
)

# Lưu lại tên input chính (lấy input đầu tiên)
inputs = self.ort_session.get_inputs()
if len(inputs) == 0:
    raise RuntimeError("ONNX model has no inputs.")
self.input_name = inputs[0].name

if verbose:
    print(f"[ZIM_Encoder] Using input name: {self.input_name}")
    print("[ZIM_Encoder] Model outputs:")
    for i, out in enumerate(self.ort_session.get_outputs()):
        print(f"  #{i}: name={out.name}, shape={out.shape},
type={out.type}")

def forward(
    self,
    image: torch.Tensor,
) -> Tuple[torch.Tensor, Tuple[torch.Tensor, torch.Tensor, torch.Tensor]]:
    """
    image: torch.Tensor shape (B, 3, 1024, 1024) trên CPU hoặc GPU
    returns: (image_embeddings, (feat_D0, feat_D1, feat_D2))

```

Lưu ý:

- Hàm này giả định model ONNX trả về 4 output:
  - [0] image\_embeddings
  - [1] feat\_D0
  - [2] feat\_D1

```

[3] feat_D2
"""

if not isinstance(image, torch.Tensor):
    raise TypeError("image must be a torch.Tensor")

if image.ndim != 4 or image.shape[1] != 3 or image.shape[2] != 1024 or
image.shape[3] != 1024:
    raise ValueError(
        f"Expected image shape (B, 3, 1024, 1024) but got
{tuple(image.shape)}. "
        "Hãy chắc chắn đã resize/normalize đúng trước khi gọi encoder."
    )

device = image.device

# Chuẩn bị numpy input (luôn về CPU)
image_np = tensor2np(image)
ort_inputs = {self.input_name: image_np}

# Chạy ONNX inference
outputs = self.ort_session.run(None, ort_inputs)

if len(outputs) != 4:
    raise RuntimeError(
        f"Expected 4 outputs (image_embeddings, feat_D0, feat_D1, feat_D2)
but got {len(outputs)}. "
        "Kiểm tra lại model ONNX hoặc chỉnh lại wrapper cho đúng số
output."
    )

image_embeddings_np, feat_D0_np, feat_D1_np, feat_D2_np = outputs

# Convert ngược về torch.Tensor trên device ban đầu
image_embeddings = np2tensor(image_embeddings_np, device)
feat_D0 = np2tensor(feat_D0_np, device)
feat_D1 = np2tensor(feat_D1_np, device)
feat_D2 = np2tensor(feat_D2_np, device)

```

```

        return image_embeddings, (feat_D0, feat_D1, feat_D2)

    def __call__(self, *args, **kwargs) -> Any:
        return self.forward(*args, **kwargs)

```

Bước 3: Tạo file mapper.py cùng thư mục với model và zim\_encoder.py

```

#!/usr/bin/env python3
# -*- coding: utf-8 -*-

import os
import sys
import io
import tarfile
import shutil
import tempfile
import subprocess
import traceback
import gc
import numpy as np
import torch
from PIL import Image
from zim_encoder import ZIM_Encoder

# ===== CẤU HÌNH MÔI TRƯỜNG =====
extra_paths = [
    r"C:\Users\lonovo\AppData\Local\Programs\Python\Python311\Lib\site-packages",
    r"C:\Users\lonovo\AppData\Roaming\Python\Python311\site-packages",
    r"C:\PythonLibs"
]
for p in extra_paths:
    if os.path.exists(p) and p not in sys.path:
        sys.path.append(p)

HADOOP_CMD = os.path.join(os.environ.get("HADOOP_HOME", r"C:\hadoop-3.3.0"),
"bin", "hadoop.cmd")
HDFS_TAR_DIR = "/hadoop/image_shards"
HDFS_NPZ_OUT_DIR = "/hadoop/zim_results_npz"

```

```

ONNX_FILE = "zim_model.onnx"
BATCH_SIZE = 4 # Tăng tốc độ bằng cách xử lý nhiều ảnh cùng lúc

def hadoop_fs(*args):
    return subprocess.check_call([HADOOP_CMD, "fs", *args], shell=True)

def preprocess_image(img_bytes):
    """Tiền xử lý ảnh đơn lẻ trả về numpy array"""
    img = Image.open(io.BytesIO(img_bytes)).convert("RGB")
    if img.size != (1024, 1024):
        img = img.resize((1024, 1024))
    return np.transpose(np.asarray(img, dtype=np.float32) / 255.0, (2, 0, 1))

def main():
    if not os.path.exists(ONNX_FILE):
        sys.stderr.write(FATAL: Khong tim thay model {ONNX_FILE}\n")
        sys.exit(1)

    # Khởi tạo Encoder một lần duy nhất
    encoder = ZIM_Encoder(onnx_path=ONNX_FILE, use_cuda=torch.cuda.is_available(),
                           verbose=False)
    workdir = tempfile.mkdtemp(prefix="zim_batch_map_")

    try:
        for raw in sys.stdin:
            line = raw.strip()
            if not line: continue

            hdfs_tar = line if line.startswith("/") else
f"{HDFS_TAR_DIR}/{os.path.basename(line)}"
            shard_id = os.path.splitext(os.path.basename(hdfs_tar))[0]
            local_tar = os.path.join(workdir, f"{shard_id}.tar")
            results_dict = {}

            try:
                hadoop_fs("-get", "-f", hdfs_tar, local_tar)

                with tarfile.open(local_tar, "r:") as tf:

```

```

batch_images = []
batch_names = []

for m in tf:
    if not m.isfile() or not m.name.lower().endswith('.jpg',
'.png', '.jpeg')):
        continue

    f = tf.extractfile(m)
    if not f: continue

    try:
        # 1. Thu thập ảnh vào batch
        batch_images.append(preprocess_image(f.read()))
        batch_names.append(m.name)

        # 2. Xử lý khi batch đầy
        if len(batch_images) == BATCH_SIZE:
            input_tensor =
torch.from_numpy(np.stack(batch_images))
            emb, _ = encoder(input_tensor)

            # Lưu kết quả
            emb_np = emb.detach().cpu().numpy()
            for i, name in enumerate(batch_names):
                results_dict[name] = emb_np[i]

            # Giải phóng bộ nhớ ngay lập tức
            del input_tensor, emb, emb_np
            batch_images, batch_names = [], []
            if torch.cuda.is_available():

torch.cuda.empty_cache()

    except Exception as img_err:
        sys.stderr.write(f"Warning: Lỗi ảnh {m.name}:
{str(img_err)}\n")

    # 3. Xử lý nốt các ảnh dư thừa cuối batch

```

```

        if batch_images:
            input_tensor = torch.from_numpy(np.stack(batch_images))
            emb, _ = encoder(input_tensor)
            emb_np = emb.detach().cpu().numpy()
            for i, name in enumerate(batch_names):
                results_dict[name] = emb_np[i]
            del input_tensor, emb, emb_np

# BƯỚC 3: Lưu và Upload NPZ
if results_dict:
    local_npz = os.path.join(workdir, f"{shard_id}.npz")
    np.savez_compressed(local_npz, **results_dict)

    hdfs_npz_path = f"{HDFS_NPZ_OUT_DIR}/{shard_id}.npz"
    hadoop_fs("-mkdir", "-p", HDFS_NPZ_OUT_DIR)
    hadoop_fs("-put", "-f", local_npz, hdfs_npz_path)
    print(f"NPZ_PATH\t{hdfs_npz_path}")

except Exception:
    sys.stderr.write(f"ERROR: Shard {shard_id} that bai:
{traceback.format_exc()}\n")
finally:
    # Dọn dẹp mạnh mẽ sau mỗi Shard
    if os.path.exists(local_tar): os.remove(local_tar)
    results_dict.clear()
    gc.collect()

finally:
    shutil.rmtree(workdir, ignore_errors=True)

if __name__ == "__main__":
    main()

```

Bước 4: Tạo file reducer.py cùng thư mục với mapper.py và zim\_encoder.py

```

#!/usr/bin/env python3
# -*- coding: utf-8 -*-

```

```

import sys
import os
import io # Đã bổ sung import này
import subprocess
import numpy as np
import tempfile
import shutil
from datetime import datetime

# Ép kiểu stdout và stderr sử dụng UTF-8 để in được tiếng Việt trên Windows/Hadoop
sys.stdout = io.TextIOWrapper(sys.stdout.buffer, encoding='utf-8')
sys.stderr = io.TextIOWrapper(sys.stderr.buffer, encoding='utf-8')

# ===== CẤU HÌNH HỆ THỐNG =====
HADOOP_HOME = os.environ.get("HADOOP_HOME", r"C:\hadoop-3.3.0")
HADOOP_CMD = os.path.join(HADOOP_HOME, "bin", "hadoop.cmd")
HDFS_REPORT_DIR = "/hadoop/reports/zim_evaluation"

def hadoop_fs(*args):
    """Thực thi lệnh shell HDFS"""
    cmd = [HADOOP_CMD, "fs", *args]
    try:
        subprocess.check_call(cmd, shell=True)
    except subprocess.CalledProcessError as e:
        sys.stderr.write(f"Reducer HDFS Error: {e}\n")

def calculate_cosine_similarity(vec_a, vec_b):
    """
    Tính độ tương đồng Cosine giữa hai vector:
    
$$\text{Similarity} = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|}$$

    """
    dot_product = np.dot(vec_a, vec_b)
    norm_a = np.linalg.norm(vec_a)
    norm_b = np.linalg.norm(vec_b)
    return dot_product / (norm_a * norm_b + 1e-8)

def main():

```

```

all_embeddings = []

workdir = tempfile.mkdtemp(prefix="zim_red_final_")

try:
    # BƯỚC 1: Thu thập toàn bộ dữ liệu từ các file .npz (Mapper gửi sang)
    for line in sys.stdin:
        line = line.strip()
        if not line: continue

        try:
            # Phân tách Key-Value từ Shuffle & Sort (NPZ_PATH \t path)
            parts = line.split("\t")
            if len(parts) < 2: continue
            hdfs_npz_path = parts[1]

            local_npz = os.path.join(workdir, os.path.basename(hdfs_npz_path))
            hadoop_fs("-get", "-f", hdfs_npz_path, local_npz)

            with np.load(local_npz) as data:
                for img_name in data.files:
                    all_embeddings.append((img_name,
data[img_name].flatten()))

            if os.path.exists(local_npz): os.remove(local_npz)
        except Exception as e:
            sys.stderr.write(f"Lỗi khi xử lý file NPZ: {str(e)}\n")

    if not all_embeddings:
        sys.stderr.write("Không tìm thấy dữ liệu để xử lý.\n")
        return

    # BƯỚC 2: Phân tích thống kê & Tính toán Global Centroid
    feature_matrix = np.array([item[1] for item in all_embeddings])
    global_centroid = np.mean(feature_matrix, axis=0)

    total_mean = np.mean(feature_matrix)
    total_std = np.std(feature_matrix)
    threshold = 0.7 # Ngưỡng tin cậy đánh giá đặc trưng ViT-B

```

```

# BƯỚC 3: Xây dựng nội dung báo cáo
report = []
report.append("-"*80)
report.append(f"HỆ THỐNG ĐÁNH GIÁ ĐỘ TIN CẬY IMAGE ENCODER (ViT-B)")
report.append(f"Thời gian tạo báo cáo:")
{datetime.now().strftime('%Y-%m-%d %H:%M:%S')}")
report.append("-"*80)

reliable_count = 0
for img_name, emb_vector in all_embeddings:
    score = calculate_cosine_similarity(emb_vector, global_centroid)
    status = "OK" if score >= threshold else "LOW_CONFIDENCE"
    if score >= threshold: reliable_count += 1
    report.append(f"{img_name}<50} | {score:.4f} [{status}]")

summary = [
    "-"*80,
    "TỔNG KẾT THỐNG KÊ:",
    f"- Tổng số lượng ảnh: {len(all_embeddings)}",
    f"- Giá trị trung bình Embedding (Mean): {total_mean:.6f}",
    f"- Độ lệch chuẩn Embedding (Std): {total_std:.6f}",
    f"- Số lượng ảnh đạt chuẩn (>{threshold}): {reliable_count}",
    f"- Tỷ lệ tin cậy toàn hệ thống:
{(reliable_count/len(all_embeddings))*100:.2f}%",
    "="*80
]
report.extend(summary)

# BƯỚC 4: Lưu báo cáo vào HDFS
final_report_str = "\n".join(report)
local_report_file = os.path.join(workdir, "zim_reliability_report.txt")

with open(local_report_file, "w", encoding="utf-8") as f:
    f.write(final_report_str)

```

```

        hadoop_fs("-mkdir", "-p", HDFS_REPORT_DIR)
        hadoop_fs("-put", "-f", local_report_file,
f"{HDFS_REPORT_DIR}/reliability_report.txt")


        # BƯỚC 5: Xuất ra stdout để ghi vào thư mục output của Job
        print(final_report_str)

    finally:
        shutil.rmtree(workdir, ignore_errors=True)

if __name__ == "__main__":
    main()

```

## Bước 5: Di chuyển đến thư mục chứa các file code



## Bước 6: Chạy chương trình

```


hadoop jar "C:/hadoop-3.3.0/share/hadoop/tools/lib/hadoop-streaming-3.3.0.jar" ^
-D mapreduce.job.name="ZIM_ViT_B_Final" ^
-D mapreduce.job.reduces=1 ^
-D mapreduce.map.memory.mb=7168 ^
-D mapreduce.map.java.opts="-Xmx5120m" ^

```

```

-D mapreduce.reduce.memory.mb=7168 ^
-D mapreduce.reduce.java.opts="-Xmx5120m" ^
-D mapreduce.task.timeout=3600000 ^
-D
mapreduce.map.env="PYTHONPATH=C:\Users\lonovo\AppData\Local\Programs\Python\Python311\Lib\site-packages;C:\PythonLibs" ^
-D
mapreduce.reduce.env="PYTHONPATH=C:\Users\lonovo\AppData\Local\Programs\Python\Python311\Lib\site-packages;C:\PythonLibs" ^
-files "mapper.py,reducer.py,zim_encoder.py,zim_model.onnx" ^
-input "/hadoop/input_shard_manifest/manifest_shards.txt" ^
-output "/hadoop/zim_feature_output/result_v1" ^
-mapper
"C:/Users/lonovo/AppData/Local/Programs/Python/Python311/python.exe
mapper.py" ^
-reducer
"C:/Users/lonovo/AppData/Local/Programs/Python/Python311/python.exe
reducer.py"

```



The screenshot shows a Microsoft Word document titled "MapReduce.docx". The content of the document is a command-line script (Administrator: Command Prompt) for running a Hadoop job. The script includes various Hadoop configuration parameters like memory and timeout, and specifies the mapper and reducer scripts to use Python files.

```

Administrator: Command Prompt
Microsoft Windows [Version 10.0.26100.7462]
(c) Microsoft Corporation. All rights reserved.

C:\Windows\System32>start-all
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons

C:\Windows\System32>cd /d D:\Hoc_ki_7\Phan_tich_du_lieu_lon\Code\ZIM\mapreducehadoop jar "C:/hadoop-3.3.0/share/hadoop/tools/lib/hadoop-streaming-3.3.0.jar" ^
More? -D mapreduce.job.name="ZIM_VIT_B_Final" ^
More? -D mapreduce.job.reduces=1 ^
More? -D mapreduce.map.memory.mb=7168 ^
More? -D mapreduce.map.java.opts="-Xmx5120m" ^
More? -D mapreduce.reduce.memory.mb=7168 ^
More? -D mapreduce.reduce.java.opts="-Xmx5120m" ^
More? -D mapreduce.task.timeout=3600000 ^
More? -D mapreduce.map.env="PYTHONPATH=C:\Users\lonovo\AppData\Local\Programs\Python\Python311\Lib\site-packages;C:\PythonLibs" ^
More? -D mapreduce.reduce.env="PYTHONPATH=C:\Users\lonovo\AppData\Local\Programs\Python\Python311\Lib\site-packages;C:\PythonLibs" ^
More? -files "mapper.py,reducer.py,zim_encoder.py,zim_model.onnx" ^
More? -input "/hadoop/input_shard_manifest/manifest_shards.txt" ^
More? -output "/hadoop/zim_feature_output/result_v1" ^
More? -mapper "C:/Users/lonovo/AppData/Local/Programs/Python/Python311/python.exe mapper.py" ^
More? -reducer "C:/Users/lonovo/AppData/Local/Programs/Python/Python311/python.exe reducer.py"

```

Bước 7: Kiểm tra kết quả thu được

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

## Browse Directory

/hadoop									Go!				
Show 25 entries									Search:				
	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name					
□	drwxr-xr-x	lonovo	supergroup	0 B	Dec 27 21:08	0	0 B	image_shards					
□	drwxr-xr-x	lonovo	supergroup	0 B	Dec 27 18:40	0	0 B	input_shard_manifest					
□	drwxr-xr-x	lonovo	supergroup	0 B	Dec 27 22:33	0	0 B	reports					
□	drwxr-xr-x	lonovo	supergroup	0 B	Dec 27 21:54	0	0 B	zim_feature_output					
□	drwxr-xr-x	lonovo	supergroup	0 B	Dec 27 22:32	0	0 B	zim_results_npz					

Showing 1 to 5 of 5 entries

Previous **1** Next

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

## Browse Directory

/hadoop/zim_results_npz									Go!				
Show 25 entries									Search:				
	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name					
□	-rw-r--r--	lonovo	supergroup	36.24 MB	Dec 27 21:56	1	128 MB	shard_01.npz					
□	-rw-r--r--	lonovo	supergroup	36.21 MB	Dec 27 21:58	1	128 MB	shard_02.npz					
□	-rw-r--r--	lonovo	supergroup	36.19 MB	Dec 27 22:00	1	128 MB	shard_03.npz					
□	-rw-r--r--	lonovo	supergroup	36.21 MB	Dec 27 22:02	1	128 MB	shard_04.npz					
□	-rw-r--r--	lonovo	supergroup	36.17 MB	Dec 27 22:04	1	128 MB	shard_05.npz					
□	-rw-r--r--	lonovo	supergroup	36.18 MB	Dec 27 22:06	1	128 MB	shard_06.npz					
□	-rw-r--r--	lonovo	supergroup	36.09 MB	Dec 27 22:08	1	128 MB	shard_07.npz					
□	-rw-r--r--	lonovo	supergroup	36.11 MB	Dec 27 22:10	1	128 MB	shard_08.npz					
□	-rw-r--r--	lonovo	supergroup	36.22 MB	Dec 27 22:11	1	128 MB	shard_09.npz					
□	-rw-r--r--	lonovo	supergroup	36.2 MB	Dec 27 22:13	1	128 MB	shard_10.npz					
□	-rw-r--r--	lonovo	supergroup	36.17 MB	Dec 27 22:15	1	128 MB	shard_11.npz					
□	-rw-r--r--	lonovo	supergroup	36.19 MB	Dec 27 22:17	1	128 MB	shard_12.npz					
□	-rw-r--r--	lonovo	supergroup	36.17 MB	Dec 27 22:19	1	128 MB	shard_13.npz					

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

## Browse Directory

/hadoop/zim_feature_output/result_v1									Go!				
Show 25 entries									Search:				
	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name					
□	-rw-r--r--	lonovo	supergroup	0 B	Dec 27 22:33	1	128 MB	_SUCCESS					
□	-rw-r--r--	lonovo	supergroup	13.78 KB	Dec 27 22:33	1	128 MB	part-00000					

Showing 1 to 2 of 2 entries

Previous **1** Next

Hadoop, 2020.

Hadoop Overview Datanodes DataNode Volume Failures Snapshot Start/Stop Utilities

## Browse Directory

/hadoop/reports/zim\_evaluation

Show 25 entries

	Permission	Owner
	-rw-r--r--	lenovo

Showing 1 to 1 of 1 entries

Hadoop, 2020.

File information - reliability\_report.txt

Download Head the file (first 32K) Tail the file (last 32K)

Block information - Block 0

Block ID: 1073743965  
Block Pool ID: BP-1268425478-192.168.56.1-1763953109125  
Generation Stamp: 3150  
Size: 14110  
Availability:

- 192.168.1.5

File contents

```
TỔNG KẾT THỐNG KẾ:
- Tổng số lượng ảnh: 200
- Giá trị trung bình Embedding (Mean): 0.013929
- Độ lệch chuẩn Embedding (Std): 0.149071
- Số lượng ảnh đạt chuẩn (>0.7): 196
- Tỷ lệ tin cậy toàn hệ thống: 98.00%
```

Close

## Phần 2: Chạy Map-Reduce với Prompt Encoder + Transformer Decoder

### Bước 1: Chuẩn bị dữ liệu

#### 1. Dữ liệu dạng RGB Image (.jpg)



#### 2. Dữ liệu của các Point, Box

```

1   "points": [[810, 540]], "point_labels": [1], "box": [405, 270, 1215, 810], "image": "hdfs:///zim_dat
2   "points": [[540, 810]], "point_labels": [1], "box": [270, 405, 810, 1215], "image": "hdfs:///zim_dat
3   "points": [[540, 570]], "point_labels": [1], "box": [270, 285, 810, 855], "image": "hdfs:///zim_data
4   "points": [[540, 810]], "point_labels": [1], "box": [270, 405, 810, 1215], "image": "hdfs:///zim_dat
5   "points": [[815, 540]], "point_labels": [1], "box": [408, 270, 1223, 810], "image": "hdfs:///zim_dat
6   "points": [[540, 810]], "point_labels": [1], "box": [270, 405, 810, 1215], "image": "hdfs:///zim_dat
7   "points": [[814, 540]], "point_labels": [1], "box": [407, 270, 1222, 810], "image": "hdfs:///zim_dat
8   "points": [[540, 810]], "point_labels": [1], "box": [270, 405, 810, 1215], "image": "hdfs:///zim_dat
9   "points": [[959, 540]], "point_labels": [1], "box": [480, 270, 1439, 810], "image": "hdfs:///zim_dat
10  "points": [[810, 540]], "point_labels": [1], "box": [405, 270, 1215, 810], "image": "hdfs:///zim_dat
11  "points": [[808, 540]], "point_labels": [1], "box": [404, 270, 1212, 810], "image": "hdfs:///zim_dat
12  "points": [[809, 540]], "point_labels": [1], "box": [405, 270, 1214, 810], "image": "hdfs:///zim_dat
13  "points": [[810, 540]], "point_labels": [1], "box": [405, 270, 1215, 810], "image": "hdfs:///zim_dat
14  "points": [[810, 540]], "point_labels": [1], "box": [405, 270, 1215, 810], "image": "hdfs:///zim_dat
15  "points": [[540, 810]], "point_labels": [1], "box": [270, 405, 810, 1215], "image": "hdfs:///zim_dat
16  "points": [[870, 540]], "point_labels": [1], "box": [435, 270, 1306, 810], "image": "hdfs:///zim_dat
17  "points": [[540, 809]], "point_labels": [1], "box": [270, 405, 810, 1214], "image": "hdfs:///zim_dat
18  "points": [[719, 540]], "point_labels": [1], "box": [360, 270, 1079, 810], "image": "hdfs:///zim_dat
19  "points": [[810, 540]], "point_labels": [1], "box": [405, 270, 1215, 810], "image": "hdfs:///zim_dat
20  "points": [[809, 540]], "point_labels": [1], "box": [405, 270, 1214, 810], "image": "hdfs:///zim_dat
21  "points": [[810, 540]], "point_labels": [1], "box": [405, 270, 1215, 810], "image": "hdfs:///zim_dat
22  "points": [[810, 540]], "point_labels": [1], "box": [405, 270, 1215, 810], "image": "hdfs:///zim_dat
23  "points": [[540, 810]], "point_labels": [1], "box": [270, 405, 810, 1215], "image": "hdfs:///zim_dat
24  "points": [[959, 540]], "point_labels": [1], "box": [480, 270, 1439, 810], "image": "hdfs:///zim_dat
25  "points": [[810, 540]], "point_labels": [1], "box": [405, 270, 1215, 810], "image": "hdfs:///zim_dat
26  "points": [[540, 726]], "point_labels": [1], "box": [270, 363, 810, 1090], "image": "hdfs:///zim_dat
27  "points": [[810, 540]], "point_labels": [1], "box": [405, 270, 1215, 810], "image": "hdfs:///zim_dat
28  "points": [[810, 540]], "point_labels": [1], "box": [405, 270, 1215, 810], "image": "hdfs:///zim_dat
29  "points": [[810, 540]], "point_labels": [1], "box": [405, 270, 1215, 810], "image": "hdfs:///zim_dat
30  "points": [[810, 540]], "point_labels": [1], "box": [405, 270, 1215, 810], "image": "hdfs:///zim_dat

```

## Bước 2: Tạo thư mục chứa các file dữ liệu trên hdfs


**hdfs dfs -mkdir -p /zim\_dataset/images**

**hdfs dfs -mkdir -p /zim\_dataset/ prompts**

Name	Type	Size	Last Modified	Replication	Block Size	Owner	Group	Permission
_tmp_test	File	0 B	Dec 27 00:06	0	0 B	hoang	supergroup	drwxr-xr-x
dataset_2	File	0 B	Dec 28 20:41	0	0 B	hoang	supergroup	drwxr-xr-x
images	File	0 B	Dec 21 22:54	0	0 B	hoang	supergroup	drwxr-xr-x
mr_eval_out	File	0 B	Dec 27 23:43	0	0 B	hoang	supergroup	drwxr-xr-x
output_masks_run1	File	0 B	Dec 27 10:07	0	0 B	hoang	supergroup	drwxr-xr-x
prompts	File	0 B	Dec 26 23:29	0	0 B	hoang	supergroup	drwxr-xr-x

### Bước 3: Đẩy dữ liệu gồm các ảnh và các Prompt chứa tọa độ các Point, box lên HDFS

```
hdfs dfs -put C:\Users\hoang\Downloads\TestMapreduce\image\*  
/zim_dataset/images/
```



```
hdfs dfs -put prompts_hdfs.jsonl /zim_dataset/prompts/
```

### Browse Directory

Name	Size	Last Modified	Replication	Block Size
prompts_hdfs.jsonl	22.95 KB	Dec 26 23:29	1	128 MB

Show 25 entries

Search:


Showing 1 to 1 of 1 entries

Previous **1** Next

Hadoop, 2020.

## Bước 4: Chuẩn bị model và code để chạy mapreduce

- Tại model encoder từ link: [https://huggingface.co/naver-iv/zim-anything-vitb/tree/main/zim\\_vit\\_b\\_2043](https://huggingface.co/naver-iv/zim-anything-vitb/tree/main/zim_vit_b_2043)



- Xây dựng code Prompt\_Encoder có tên là prompt\_stage.py

```
from __future__ import annotations
from dataclasses import dataclass
from typing import Optional, Tuple

import numpy as np
import torch

@dataclass
class PromptPack:
    # Tensor format đúng với predictor/model.decoder đang dùng
    points: Optional[Tuple[torch.Tensor, torch.Tensor]] # (coords[B,N,2], labels[B,N])
    boxes: Optional[torch.Tensor] # [B,4] XYXY
    attn_mask: Optional[torch.Tensor] # [B,S,S] (nếu dùng)

class PromptStage:
    """
    'Prompt Encoder' theo nghĩa hệ thống (không phải class PromptEncoder của SAM):
        - transform point coords theo predictor.transform
        - convert sang torch tensor
        - tạo attention mask (nếu model yêu cầu)
        - trả về PromptPack cho decoder stage
    """
    def __init__(self, transform, model, device: torch.device):
        self.transform = transform # từ predictor (ResizeLongestSide hoặc tương đương)
        self.model = model # ZIM model wrapper (có point_attn_mask)
        self.device = device

    @torch.no_grad()
    def build(self):
```

```

        self,
        original_size: Tuple[int, int],
        point_coords: Optional[np.ndarray] = None,      # [N,2] pixel theo ảnh
gốc
        point_labels: Optional[np.ndarray] = None,      # [N]
        boxes: Optional[np.ndarray] = None,             # [4] hoặc [B,4] theo ảnh
gốc
) -> PromptPack:

    points_torch = None
    boxes_torch = None
    attn_mask = None

    # 1) Points
    if point_coords is not None:
        assert point_labels is not None, "point_labels must be supplied if
point_coords is supplied."

        # predictor.py đang làm apply_coords(point_coords, original_size)
        point_coords_t = self.transform.apply_coords(point_coords,
original_size)

        coords_torch = torch.as_tensor(point_coords_t,
dtype=torch.float32, device=self.device)
        labels_torch = torch.as_tensor(point_labels, dtype=torch.float32,
device=self.device)

        # predictor.py dùng shape [1,N,2] và [1,N]
        coords_torch = coords_torch[None, :, :]
        labels_torch = labels_torch[None, :]

        points_torch = (coords_torch, labels_torch)

    # 2) Attention mask (theo predictor.py:
self.model.point_attn_mask(point_coords))
        # Lưu ý: model.point_attn_mask expects point_coords [B,N,2]
        attn_mask = self.model.point_attn_mask(coords_torch)

    # 3) Boxes
    if boxes is not None:
        # hỗ trợ boxes dạng [4] hoặc [B,4]
        if boxes.ndim == 1:
            boxes = boxes[None, :]
        # transform boxes theo original_size (nếu predictor có
apply_boxes)
            if hasattr(self.transform, "apply_boxes"):
                boxes_t = self.transform.apply_boxes(boxes, original_size)
            else:
                # fallback: transform 2 điểm (x0,y0) và (x1,y1)
                # (chỉ dùng khi không có apply_boxes)
                boxes_t = boxes.copy()

```

```

        boxes_t[:, 0:2] = self.transform.apply_coords(boxes[:, 0:2],
original_size)
        boxes_t[:, 2:4] = self.transform.apply_coords(boxes[:, 2:4],
original_size)

        boxes_torch = torch.as_tensor(boxes_t, dtype=torch.float32,
device=self.device)

    return PromptPack(points=points_torch, boxes=boxes_torch,
attn_mask=attn_mask)

```

### 3. Xây dựng Transformer\_Decoder có tên là decode\_stage.py

```

import torch
import torch.nn.functional as F

class DecodeStage:
    def __init__(self, model):
        self.model = model

    @torch.no_grad()
    def run(
        self,
        image_embeddings,
        interm_feats,           # <-- BẮT BUỘC cho ZIM decoder
        points=None,
        boxes=None,
        attn_mask=None,
        multimask_output=False
    ):
        # ZIM decoder requires interm_feats
        masks, iou_predictions = self.model.decoder(
            image_embeddings=image_embeddings,
            interm_feats=interm_feats,
            points=points,
            boxes=boxes,
            attn_mask=attn_mask,
        )

        # low_res_masks: giữ tương tự predictor.py (upsample x2)
        low_res_masks = F.interpolate(
            masks,
            scale_factor=2,
            mode="bilinear",
            align_corners=False
        )

        # output dạng sigmoid như predictor

```

```
        return masks.sigmoid(), iou_predictions, low_res_masks.sigmoid()
```

4. Xây dựng Mapper để có thể chạy module với Mapreduce lấy tên là mapper\_zim.py

```
import sys
import json
import os
import tempfile
import subprocess
import base64

import numpy as np
import cv2
import torch

from zimAnything import zimModelRegistry, ZimPredictor
from zimAnything.modular.promptStage import PromptStage
from zimAnything.modular.decodeStage import DecodeStage

# ===== HDFS HELPERS =====

def hdfs_cmd() -> str:
    """
    Return path to hdfs executable.
    On Windows, prefer %HADOOP_HOME%\bin\hdfs.cmd
    """
    hadoop_home = os.environ.get("HADOOP_HOME") or
    os.environ.get("HADOOP_PREFIX")
    if hadoop_home:
        cand = os.path.join(hadoop_home, "bin", "hdfs.cmd")
        if os.path.exists(cand):
            return cand
    # Fallback (if hdfs is in PATH)
    return "hdfs"

def hdfs_get(hdfs_path: str, local_path: str) -> None:
    """
    Download a file from HDFS to local_path.
    Supports hdfs:///path style by converting to /path.
    """
    if hdfs_path.startswith("hdfs://"):
        hdfs_path = hdfs_path.replace("hdfs://", "/")
    cmd = [hdfs_cmd(), "dfs", "-get", "-f", hdfs_path, local_path]
    subprocess.check_call(cmd, stdout=subprocess.DEVNULL,
    stderr=subprocess.DEVNULL)
```

```

# ===== OUTPUT ENCODING =====

def mask_to_png_b64(mask01: np.ndarray) -> str:
    """
    mask01: float/0-1 array -> PNG bytes -> base64 ascii
    """
    m = (mask01 > 0.5).astype(np.uint8) * 255
    ok, buf = cv2.imencode(".png", m)
    if not ok:
        raise RuntimeError("encode png failed")
    return base64.b64encode(buf.tobytes()).decode("ascii")

def emit(obj: dict) -> None:
    sys.stdout.write(json.dumps(obj, ensure_ascii=False) + "\n")
    sys.stdout.flush()

# ===== MAIN =====

def main():
    # ---- checkpoint dir (required) ----
    checkpoint_dir = os.environ.get("ZIM_CKPT_DIR")
    if not checkpoint_dir:
        raise RuntimeError(
            "Missing env var ZIM_CKPT_DIR. "
            "Set it to folder containing encoder.onnx and decoder.onnx, "
            "e.g. C:\\Users\\hoang\\Downloads\\BTL-"
            "BigData\\ZIM\\results\\zim_vit_l_2092"
        )

    if not os.path.isdir(checkpoint_dir):
        raise RuntimeError(f"ZIM_CKPT_DIR is not a directory: {checkpoint_dir}")

    enc_path = os.path.join(checkpoint_dir, "encoder.onnx")
    dec_path = os.path.join(checkpoint_dir, "decoder.onnx")
    if not os.path.exists(enc_path):
        raise RuntimeError(f"encoder.onnx not found in ZIM_CKPT_DIR: {enc_path}")
    if not os.path.exists(dec_path):
        raise RuntimeError(f"decoder.onnx not found in ZIM_CKPT_DIR: {dec_path}")

    # ---- device ----
    device = "cuda" if torch.cuda.is_available() else "cpu"

    # ---- build model ----
    keys = list(zim_model_registry.keys())
    if not keys:
        raise RuntimeError("zim_model_registry empty")

```

```

model_key = keys[0] # bạn có thể hardcode key cụ thể nếu muốn

model = zim_model_registry[model_key](checkpoint=checkpoint_dir)
model.to(device).eval()

predictor = ZimPredictor(model)
prompt_stage = PromptStage(
    predictor.transform,
    predictor.model,
    next(predictor.model.parameters()).device
)

decode_stage = DecodeStage(predictor.model)

# ---- streaming loop ----
for line in sys.stdin:
    line = line.strip()
    if not line:
        continue

    try:
        rec = json.loads(line)

        rid = rec.get("id", "")
        hdfs_img = rec["image"]

        points = rec.get("points")           # [[x,y], ...] theo ảnh gốc
        point_labels = rec.get("point_labels") # [1/0, ...]
        box = rec.get("box")                 # [x1,y1,x2,y2] theo ảnh gốc

        # --- fetch image from HDFS to local temp ---
        with tempfile.TemporaryDirectory() as td:
            local_img = os.path.join(td, "img.jpg")
            hdfs_get(hdfs_img, local_img)

            bgr = cv2.imread(local_img)
            if bgr is None:
                raise RuntimeError(f"cv2.imread failed: {local_img}")
            rgb = cv2.cvtColor(bgr, cv2.COLOR_BGR2RGB)

        # --- encoder stage ---
        # --- encoder stage ---
        predictor.set_image(rgb) # tạo predictor.features + (thường)
    interm_feats
        image_embeddings = predictor.features
        interm_feats = get_interm_feats_from_predictor(predictor)

        if interm_feats is None:
            raise RuntimeError(
                "Cannot find interm_feats on predictor after set_image().

"

```

```

        "Add debug to print predictor attrs containing
'feat'/'inter'."
    )

    # --- prepare prompt inputs ---
    pc = pl = bx = None
    if points is not None and point_labels is not None:
        pc = np.array(points, dtype=np.float32)
        pl = np.array(point_labels, dtype=np.float32)
    if box is not None:
        bx = np.array(box, dtype=np.float32)

    # --- prompt stage ---
    pack = prompt_stage.build(
        original_size=predictor.original_size,
        point_coords=pc,
        point_labels=pl,
        boxes=bx,
    )

    # --- decode stage ---
    masks, iou, low_res = decode_stage.run(
        image_embeddings=image_embeddings,
        interm_feats=interm_feats,
        points=pack.points,
        boxes=pack.boxes,
        attn_mask=pack.attn_mask,
        multimask_output=False,
    )

    mask01 = masks[0, 0].detach().cpu().numpy()
    iou_score = float(iou[0, 0].detach().cpu().numpy())

    emit({
        "id": rid,
        "image": hdfs_img,
        "iou": iou_score,
        "mask_png_b64": mask_to_png_b64(mask01),
    })

except Exception as e:
    emit({
        "id": rec.get("id", "") if isinstance(locals().get("rec"),
dict) else "",
        "error": str(e),
        "line": line[:300],
    })

```

```

def get_interm_feats_from_predictor(predictor):
    """
    Try to retrieve intermediate features produced by encoder.
    ZIM decoder requires these.
    """
    for name in ["interm_feats", "interm_features", "features_interm", "feat",
    "feats", "interm"]:
        if hasattr(predictor, name):
            val = getattr(predictor, name)
            if val is not None:
                return val
    return None

if __name__ == "__main__":
    main()

```

## 5. Tạo file Reducer lấy tên là reducer\_avg\_iou.py

```

import sys, json

def main():
    sum_iou = 0.0
    count = 0
    min_iou = None
    max_iou = None

    # Streaming reducer input thường là: key \t value
    for line in sys.stdin:
        line = line.strip()
        if not line:
            continue

        try:
            # hỗ trợ cả 2 kiểu:
            # 1) "ALL\t{json...}"
            # 2) "{json...}" (phòng khi mapper không emit key)
            if "\t" in line:
                _, v = line.split("\t", 1)
            else:
                v = line

            obj = json.loads(v)
            iou = float(obj.get("iou", 0.0))

            sum_iou += iou
            count += 1
            min_iou = iou if min_iou is None else min(min_iou, iou)
            max_iou = iou if max_iou is None else max(max_iou, iou)

        except Exception as e:
            print(f"Error processing line: {line}. Reason: {e}")

```

```

        except Exception:
            # bỏ qua dòng lỗi
            continue

    avg = (sum_iou / count) if count else 0.0

    out = {
        "avg_iou": avg,
        "num_samples": count,
        "min_iou": min_iou,
        "max_iou": max_iou
    }

    # reducer output: 1 dòng JSON
    sys.stdout.write(json.dumps(out, ensure_ascii=False) + "\n")

if __name__ == "__main__":
    main()

```

## Bước 5: Di chuyển đến thư mục chứa code để chuẩn bị chạy Map-Reduce

```
cd "C:\Users\hoang\Downloads\BTL-BigData\ZIM"
```

## Bước 6: Thực hiện chạy Map-Reduce

```

# Set env cho phiên PowerShell hiện tại

$env:ZIM_CKPT_DIR = "C:/Users/hoang/Downloads/BTL-
BigData/ZIM/results/zim_vit_l_2092"

# Xoá output cũ

Remove-Item -Recurse -Force .\mr_eval_out -ErrorAction
SilentlyContinue

# Đường dẫn hadoop + jar

$hadoop = Join-Path $env:HADOOP_HOME "bin\hadoop.cmd"

$jar = Join-Path $env:HADOOP_HOME
"share\hadoop\tools\lib\hadoop-streaming-3.3.0.jar"

# Chạy streaming local (MAP + REDUCE)

& $hadoop jar $jar `
```

```

-D mapreduce.framework.name=local `

-fs file:/// `

-input file:///C:/Users/hoang/Downloads/BTL-
BigData/ZIM/prompts_local.jsonl `

-output file:///C:/Users/hoang/Downloads/BTL-BigData/ZIM/mr_eval_out
` 

-mapper "python mapper_zim.py" `

-reducer "python reducer_avg_iou.py" `


-file mapper_zim.py `

-file reducer_avg_iou.py `

-verbose

```

## Bước 7: Kết quả thu được



## Browse Directory

/zim_dataset/mr_eval_out								Go!					
Show 25 entries									Search:				
	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name					
<input type="checkbox"/>	-rw-r--r--	hoang	supergroup	0 B	Dec 27 23:43	1	128 MB	_SUCCESS					
<input type="checkbox"/>	-rw-r--r--	hoang	supergroup	116 B	Dec 27 23:43	1	128 MB	part-00000					

Showing 1 to 2 of 2 entries

Previous 1 Next

localhost:9870/explorer.html#/zim\_dataset/mr\_eval\_out

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities

File information - part-00000

Download Head the file (first 32K) Tail the file (last 32K)

Block information -- Block 0

Block ID: 1073742051  
Block Pool ID: BP-1019690126-192.168.56.1-1764554986257  
Generation Stamp: 1227  
Size: 116  
Availability:

- HoangAnhVu

File contents

```
[{"avg_iou": 0.5415910649299621, "num_samples": 100, "min_iou": 0.33850938081741333, "max_iou": 0.8379999399185181}]
```